

# Implementation Assignment 1

Deqing Qu (qud@oregonstate.edu) (50%)

Kejo Hsieh (hsiehke@oregonstate.edu) (50%)

## Part 0 Preprocessing and simple analysis

### **a) Remove the ID feature. Why do you think it is a bad idea to use this feature in learning?**

The id is not a factor that affects the price of a real estate, hence it does not include useful information towards generating a prediction. Using this set of data as a feature would cause adverse effects to our model, eventually generate bad predictions which increases the loss.

### **b) Split the date feature into three separate numerical features: month, day, and year. Can you think of better ways of using this date feature?**

Splitting the date feature into 3 parts would let us make predictions based on the time cycle within a year or a month. It might include certain trends if there are certain periods in a year or month where you can sell your house at a higher or lower price. If we only use the date feature as a whole, it only provides information based on a one-dimensional timeline, which would hardly be useful since time could not be re-winded.

### **c) Build a table that reports the statistics for each feature.**

Training Data:

(Please see next page)

Feature	Mean	Standard Dev	Range	Min	Max	Feature	Category	%
month	6.5924	3.11127984	11	1	12	Waterfront	1	0.007
day	15.8021	8.62133027	30	1	31		0	0.993
year	2014.3185	0.46589457	1	2014	2015	Condition	1	0.001
bedrooms	3.3752	0.94319932	32	1	33		2	0.008
bathrooms	2.118875	0.76508985	7.25	0.5	7.75		3	0.653
sqft_living	2080.2232	911.28879	9520	370	9890		4	0.257
sqft_lot	15089.2014	41201.8347	1650787	572	1651359		5	0.081
floors	1.5037	0.54261986	2.5	1	3.5	Grade	1	0
waterfront	0.007	0.08337266	1	0	1		2	0
view	0.2294	0.75589393	4	0	4		3	0
condition	3.4091	0.65355733	4	1	5		4	0.001
grade	7.6732	1.18000075	9	4	13		5	0.011
sqft_above	1793.0993	830.82389	8490	370	8860		6	0.093
sqft_basement	287.1239	434.983513	2720	0	2720		7	0.413
yr_built	1971.1249	29.4791197	115	1900	2015		8	0.284
yr_renovated	1973.4386	28.8820261	115	1900	2015		9	0.118
zipcode	0.4146	0.49265286	1	0	1		10	0.055
lat	47.5598142	0.13864366	0.6217	47.1559	47.7776		11	0.021
long	-122.21329	0.14139761	1.195	-122.514	-121.319		12	0.004
sqft_living15	1994.3261	691.865705	5650	460	6110	ZIP Code	13	0.001
sqft_lot15	12746.3234	28239.8309	870540	660	871200		980	0.585
							981	0.415

Dev:

Feature	Mean	Standard Dev	Range	Min	Max	Feature	Category	%
dummy	1	0	0	1	1	Waterfront	1	0.008
month	6.53082008	3.13310043	11	1	12		0	0.992
day	15.6751831	8.64836758	30	1	31	Condition	1	0.001
year	2014.33357	0.47148868	1	2014	2015		2	0.009
bedrooms	3.36591031	0.90524998	8	1	9		3	0.651
bathrooms	2.1113543	0.76355723	7.25	0.75	8		4	0.26
sqft_living	2073.00125	906.762023	13150	390	13540		5	0.078
sqft_lot	14601.3763	38419.1154	1023459	609	1024068	Grade	1	0
floors	1.48731463	0.53684993	2.5	1	3.5		2	0
waterfront	0.00786135	0.08831508	1	0	1		3	0
view	0.22994461	0.76634839	4	0	4		4	0.002
condition	3.40414508	0.65124828	4	1	5		5	0.013
grade	7.6482044	1.15327984	10	3	13		6	0.091
sqft_above	1784.97231	817.001899	9020	390	9410		7	0.412
sqft_basement	288.028944	441.919858	4820	0	4820		8	0.288
yr_built	1971.06754	29.1700469	115	1900	2015		9	0.125
yr_renovated	1973.52635	28.6124579	115	1900	2015		10	0.05
zipcode	0.41468644	0.49266784	1	0	1		11	0.014
lat	47.5605231	0.13902324	0.6154	47.1622	47.7776		12	0.004
long	-122.21362	0.1409749	1.196	-122.511	-121.315		13	0
sqft_living15	1977.85939	669.858563	5540	670	6210	ZIP Code	980	0.585
sqft_lot15	12812.61	27159.8441	434077	651	434728		981	0.415

Test:

Feature	Mean	Standard Dev	Range	Min	Max
month	6.5835	3.10381933	11	1	12
day	15.5078333	8.64102262	30	1	31
year	2014.3205	0.46666878	1	2014	2015
bedrooms	3.37666667	0.91712716	9	1	10
bathrooms	2.11491667	0.78033272	7.5	0.5	8
sqft_living	2087.31517	939.561862	11670	380	12050
sqft_lot	15581.0023	44341.9598	1164274	520	1164794
floors	1.48441667	0.53708984	2.5	1	3.5
waterfront	0.00816667	0.08999985	1	0	1
view	0.2465	0.78341416	4	0	4
condition	3.41633333	0.64472208	4	1	5
grade	7.6415	1.17982107	9	4	13
sqft_above	1784.47383	832.492917	8190	380	8570
sqft_basement	302.841333	455.640113	3500	0	3500
yr_built	1970.72767	29.3882942	115	1900	2015
yr_renovated	1973.15417	28.8648702	115	1900	2015
zipcode	0.41766667	0.49317464	1	0	1
lat	47.5601564	0.13794403	0.6182	47.1593	47.7775
long	-122.21548	0.13933273	1.204	-122.519	-121.315
sqft_living15	1981.94983	688.09334	5391	399	5790
sqft_lot15	12727.5395	25695.6199	411212	750	411962

Feature	Category	%
Waterfront	1	0.80%
	0	99.20%
Condition	1	0.10%
	2	0.70%
	3	64.10%
	4	27.50%
	5	7.50%
Grade	1	0.00%
	2	0.00%
	3	0.00%
	4	0.10%
	5	1.10%
	6	9.90%
	7	42.30%
	8	26.90%
	9	12.20%
	10	5.10%
	11	1.80%
	12	0.40%
	13	0.10%
ZIP Code	980	58.20%
	981	41.80%

**d ) Based on the meaning of the features as well as the statistics, which set of features do you expect to be useful for this task? Why?**

All of the features should be useful for the task. Yet we imagine the house area would be most important of them all since how big of a house should be the primary key towards its price. The year built should also be important since a house's price will depreciate with time. Overall, the grade that has received for the house should give us a comprehensive view of the price.


**e ) Normalize all features to the range between 0 and 1 using the training data.**

For each column (feature) in the data, we normalize them by deducting the smaller value in the column, then divided by the range of the column. Code implementation could be found in the `normalize_matrix()` function from our source code.

$$normalized\ x_i = \frac{x_i - x_{min}}{x_{max} - x_{min}}$$

## Part 1 Explore different learning rate

a) Which learning rate or learning rates did you observe to be good for this particular dataset? What learning rates make the gradient descent explode?



According to the result, the gradient descent will only work starting from the learning rate =  $10^{-5}$ .

Although the learning rate has to be small enough to make our gradient descent functional, it does not provide a better result when the learning rate is too small. Once it is smaller than  $10^{-5}$ , it takes a lot more iterations to get a lower norm of the gradient, hence making it harder to converge. This indicates that we are taking too small of a step towards the ideal loss range, and it would take forever to reach the goal if we don't make bigger steps.

b) For each learning rate worked for you, Report the SSE on the training data and the validation data respectively and the number of iterations needed to achieve the convergence condition for training.

What do you observe?

Learning Rate	0.00001	0.000001	0.0000001
SSE on the training set	39,370.91	39,374.84	39,491.47
SSE on the validation set	22,939.45	23,109.93	23,167.24
The number of iterations to convergence	148,433 (convergence, norm of gradient = 0.5)	500,000 (no convergence, norm of gradient = 2.90)	500,000 (no convergence, norm of gradient = 51.16)

The validation SSE is nearly half the training SSE, due to the validation dataset itself is half the size of the training dataset. If we divide the SSE for the training set in half, the SSE would be smaller than the validation set. Also, SSE increases in the training set while learning rate decreases due to the reason mentioned in 1a.

**c) Which features are the most important in deciding the house prices according to the learned weights?**

Feature	Weight	Feature	Weight
Bias	-2.80085068	Condition	1.42377189
Month	0.19954502	Grade	8.9102138
Day	-0.16887045	Sqft_above	7.91890032
Year	0.38426361	Sqft_basement	1.29948111
Bedrooms	-2.92872561	Year_built	-3.273957
Bathrooms	3.52862237	Year_Renovated	0.50622065
Sqft_living	7.20665111	Zip code	0.03393697
Sqft_lot	1.36819341	Latitude	3.51936519
Floors	-0.04122866	Longitude	-1.55944545
Waterfront	4.72380737	Sqft_living15	1.62366308
View	2.20282349	Sqft_lot15	-2.448195

(Top 3 weights are labeled in red)

According to the data, the most important feature among all the the grade that the house received, following by the house area excluding the basement, and the living room's area. These are totally understandable and matches our own assumptions. What turns out to be a surprise is that the year built was not as important as we thought. This could indicate that no matter how old the house is, its value could remain given proper management.

## Part 2 Experiments with different $\lambda$ values.

$\lambda$	SSE on the training set	SSE on the validation set
$\lambda = 0$	39370.26	23202.59
$\lambda = 0.001$	39370.27	23194.98
$\lambda = 0.01$	39370.29	23177.12
$\lambda = 0.1$	39370.58	23168.81
$\lambda = 1$	39378.37	23122.98
$\lambda = 10$	39612.20	23218.55
$\lambda = 100$	44149.37	26532.42

### Weights ( $\lambda = 0$ )

[-2.80744024 0.20010598 -0.16915518 0.38436488 -2.93083008 3.52879205  
7.45364593 2.4598755 -0.04402851 4.72516218 2.19790808 1.42664734  
8.90594427 7.69136399 1.22871644 -3.27261545 0.50914146 0.03576403  
3.5212088 -1.5537145 1.63921798 -3.23349637]

### Weights ( $\lambda = 0.001$ )

[-2.80744281 0.20010386 -0.16915492 0.38436305 -2.93070049 3.52879608  
7.40334111 2.45822904 -0.04398909 4.72509264 2.19793329 1.42663844  
8.9059034 7.73594512 1.24304808 -3.27258213 0.50911311 0.03576635  
3.52120154 -1.55370628 1.63932176 -3.23212022]

### Weights ( $\lambda = 0.01$ )

[-2.80746651 0.20008478 -0.16915263 0.38434658 -2.92953476 3.52883222  
7.2888702 2.44351929 -0.04363462 4.72446705 2.19815974 1.42655866  
8.90553509 7.83550184 1.27538371 -3.27228212 0.50885835 0.03578746  
3.52113642 -1.55363172 1.64025682 -3.21981389]

**Weights (  $\lambda = 0.1$  )**

[-2.80774213 0.19989773 -0.16913186 0.38418244 -2.91792941 3.52917922  
7.28849984 2.30446713 -0.04011816 4.71822886 2.20038677 1.42578234  
8.90181439 7.81059356 1.2718046 -3.26927543 0.50633684 0.03601053  
3.52049944 -1.55284181 1.64967672 -3.10265423]

**Weights (  $\lambda = 1$  )**

[-2.81268816 0.198260504 -0.169059267 0.382578128 -2.80644432 3.53121529  
7.24737848 1.43170374 -0.0704034182 4.65726577 2.22020923 1.41942656  
8.86199888 7.60186343 1.24854715e -3.23890102 0.482960239 0.00389757159  
3.51507115 -1.54196192 1.74684342 -2.31540946]

**Weights (  $\lambda = 10$  )**

[-2.8447067 0.18453259 -0.17011881 0.36719183 -1.96667913 3.46801241  
6.30847794 0.29895814 0.24593572 4.13002635 2.36908871 1.37956592  
8.41691333 6.65358987 1.2944505 -2.95797763 0.30663033 0.07125641  
3.47837876 -1.38370409 2.50580208 -0.65170268]

**Weights (  $\lambda = 100$  )**

[-2.11083342 0.07974424 -0.1772446 0.25931638 0.29782907 2.69297295  
4.04056444 0.17046833 1.09241682 2.05389369 2.81345462 1.10854322  
5.79047614 4.0138463 1.61390135 -1.67915477 -0.18369711 0.13630947  
3.1820594 -0.56663365 3.52228538 0.0984885 ]

**a) What trend do you observe from the training SSE as we change  $\lambda$  value?**

When  $\lambda = 0$ , it gives us the original SSE without the regularization. The SSE is at its lowest point given  $\lambda = 0$ . Then would gradually increase with the value of  $\lambda$ .

**b) What trend do you observe from the validation SSE?**

For the validation set, we can observe that giving  $\lambda \neq 0$  does decrease the SSE in a certain range. However, if the  $\lambda$  value becomes too large, the SSE will start to increase again. The best  $\lambda$  we observed in the test data set is 1.

**c) Provide an explanation for the observed behaviors.**

Since the regularize process is to prevent overfitting, the SSE for the training set will certainly increase once  $\lambda \neq 0$ , because overfitting is guaranteed to give a lower SSE for the training dataset.

For the validation data, we can see that  $\lambda$  has a best value that gives the lowest SSE, any values larger or smaller would increase the SSE. This is because  $\lambda$  determines how smooth the function is. Regularization itself is trying to control the gradient direction to be more aligned to the accumulated weight vectors we have generated through huge amounts of iterations, and  $\lambda$  determinates the amount of influence of the accumulated  $w$ . So if  $\lambda$  is too small, it would be easily effected with some rare case values. However, a  $\lambda$  that is too large would make the latter part of the data have almost no influence to the gradient descent's direction, hence increases the loss.

**d) What features get turned off for  $\lambda = 10, 10^{-2}$  and 0 ?**

$\lambda = 0$  nothing is turned off

$\lambda = 0.01$  floors

$\lambda = 10$  sqft\_lot & sqft\_lot15

If a feature gets turned off ( $|w|$  drops dramatically), it could indicate that these features are not very important on a whole scale. When a weight drops as the  $\lambda$  increase, it could indicate that these features were originally biased to seem like it's important due to some rare-case data.




## Part 3

Learning rate	SSE on the training set	SSE on the validation set
lr=1	N/A	N/A
lr=0	N/A	N/A
lr=10 <sup>-3</sup>	N/A	N/A
lr=10 <sup>-6</sup>	N/A	N/A
lr=10 <sup>-9</sup>	N/A	N/A
lr=10 <sup>-15</sup>	119580504.11	66814097.37

The SSE on the training set and validation set for learning rate = 10<sup>-15</sup>

iterations	SSE on the training set	SSE on the validation set
it=0	199375657068.76	105549427720.40
it=1000	773710507.29	494570568.08
it=2000	398835389.63	220381955.71
it=3000	304717979.74	167480855.61
it=4000	241746811.51	133157927.38
it=5000	198904631.96	109874948.41
it=6000	169674711.32	93999486.95
it=7000	149654589.58	83130551.95

it=8000	135866406.14	75647889.72
it=9000	126295879.97	70456007.79
it=10000	119580504.11	66814097.37



**a) What do you observe?**

Non-normalized data explodes very easily, it needs a smaller learning rate than the normalized data to prevent explosion.

**b) Specify the learning rate that prevents the gradient descent from exploding**

The learning rate would have to be lesser than  $10^{-15}$ .

**c) Which one is easier to train and why?**

The normalized data is much easier to train. Non-normalized data would easily explode unless the learning rate is small enough, and the weight of each feature would not be very balanced against each other due to the nature that they have very different value ranges. For example, 'year' is normally beyond 2000 but bedrooms will often be a one-digit number. Normalized data do not have this problem and is easier to control.